## Chapter One

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.
The size of the step governs convergence speed and stability. Too large a rate overshoots; too small a rate wastes iterations.

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.
The size of the step governs convergence speed and stability.
Too large a rate overshoots; too small a rate wastes iterations.

## Chapter Two

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.
The size of the step governs convergence speed and stability.
Too large a rate overshoots; too small a rate wastes iterations.

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.
The size of the step governs convergence speed and stability.
Too large a rate overshoots; too small a rate wastes iterations.