

Machine Learning

Chapter One

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.

The size of the step governs convergence speed and stability.

Too large a rate overshoots; too small a rate wastes iterations.

Machine Learning

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.

The size of the step governs convergence speed and stability.

Too large a rate overshoots; too small a rate wastes iterations.

Machine Learning

Chapter Two

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.

The size of the step governs convergence speed and stability.

Too large a rate overshoots; too small a rate wastes iterations.

Machine Learning

Gradient descent minimises the loss by stepping in the direction opposite to the gradient, scaled by the learning rate. Each step moves the parameters slightly closer to a local minimum.

The size of the step governs convergence speed and stability.

Too large a rate overshoots; too small a rate wastes iterations.