

Results and Discussion

The measurements below were collected over three runs of the training procedure, with the learning rate held constant.

- Accuracy improved on every run
- Loss decreased monotonically
- Variance across runs stayed small

Run	Accuracy
1	0.91
2	0.94

Table 1: Accuracy by run.

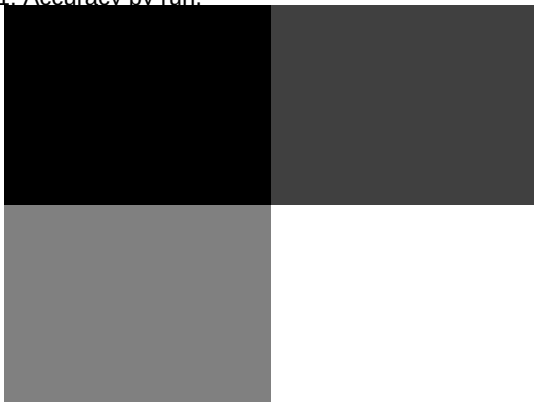


Figure 1: Loss curve over training steps.